

Examining Cars from *eBay Kleinanzeigen*

Introduction

This project examines 50,000 used car/ used car sales data points from the classifieds section of the German eBay website. The data originates from Kaggle but has been dirtied by Dataquest for purposes of data cleaning. A few of the dataset variables include:

- **name:** the name of the car
- **dateCreated:** the date the eBay listing was created
- **nrOfPictures:** the number of pictures in the ad
- **kilometer:** how many kilometers the car has driven

Here's sampling of the first few rows.

```
setwd("/Users/roberthazell/Desktop/Dataquest/Germany-Ebay-Car-Analysis")
autos <- read.csv("autos.csv")
# use auto_info to preserve original dataset auto_info
auto_info = autos
head(auto_info)
```

```
      dateCrawled
1 2016-03-26 17:47:46
2 2016-04-04 13:38:56
3 2016-03-26 18:57:24
4 2016-03-12 16:58:10
5 2016-04-01 14:38:50
6 2016-03-21 13:47:45

      name
1      Peugeot_807_160_NAVTECH_ON_BOARD
2      BMW_740i_4_4_Liter_HAMANN_UMBAU_Mega_Optik
3      Volkswagen_Golf_1.6_United
4      Smart_smart_fortwo_coupe_softouch/F1/Klima/Panorama
5      Ford_Focus_1_6_Benzin_T\xdcV_neu_ist_sehr_gepflegt.mit_Klimaanlage
6      Chrysler_Grand_Voyager_2.8_CRD_Aut.Limited_Stow\`xb4n_Go_Sitze_7Sitze
seller offerType price abtest vehicleType yearOfRegistration gearbox
1 privat Angebot $5,000 control bus 2004 manuell
2 privat Angebot $8,500 control limousine 1997 automatik
3 privat Angebot $8,990 test limousine 2009 manuell
4 privat Angebot $4,350 control kleinwagen 2007 automatik
5 privat Angebot $1,350 test kombi 2003 manuell
6 privat Angebot $7,900 test bus 2006 automatik
powerPS model odometer monthOfRegistration fuelType brand
1 158 andere 150,000km 3 lpg peugeot
2 286 7er 150,000km 6 benzin bmw
3 102 golf 70,000km 7 benzin volkswagen
4 71 fortwo 70,000km 6 benzin smart
5 0 focus 150,000km 7 benzin ford
6 150 voyager 150,000km 4 diesel chrysler
notRepairedDamage dateCreated nrOfPictures postalCode
1 nein 2016-03-26 00:00:00 0 79588
2 nein 2016-04-04 00:00:00 0 71034
```

```

3          nein 2016-03-26 00:00:00          0      35394
4          nein 2016-03-12 00:00:00          0      33729
5          nein 2016-04-01 00:00:00          0      39218
6          2016-03-21 00:00:00          0      22962
      lastSeen
1 2016-04-06 06:45:54
2 2016-04-06 14:45:08
3 2016-04-06 20:15:37
4 2016-03-15 03:16:28
5 2016-04-01 14:38:50
6 2016-04-06 09:45:21

```

Initial Data Exploration

Taking a look at the structure:

```
str(auto_info)
```

```

'data.frame':  50000 obs. of  20 variables:
 $ dateCrawled      : Factor w/ 48213 levels "2016-03-05 14:06:30",...: 31827 46100 32003 10998 41030 ...
 $ name            : Factor w/ 38754 levels "____AUDI_A4_S_LINE____VOLLAUSSTATUNG____",...: 2514 ...
 $ seller          : Factor w/ 2 levels "gewerblich","privat": 2 2 2 2 2 2 2 2 2 2 ...
 $ offerType       : Factor w/ 2 levels "Angebot","Gesuch": 1 1 1 1 1 1 1 1 1 1 ...
 $ price           : Factor w/ 2357 levels "$0","$1","$1,000",...: 1728 2181 2215 1561 58 2093 1367 1 ...
 $ abtest          : Factor w/ 2 levels "control","test": 1 1 2 1 2 2 2 1 2 1 ...
 $ vehicleType     : Factor w/ 9 levels "", "andere", "bus",...: 3 8 8 6 7 3 8 8 1 3 ...
 $ yearOfRegistration: int  2004 1997 2009 2007 2003 2006 1995 1998 2000 1997 ...
 $ gearbox         : Factor w/ 3 levels "", "automatik",...: 3 2 3 2 3 2 3 3 3 3 ...
 $ powerPS         : int  158 286 102 71 0 150 90 90 0 90 ...
 $ model           : Factor w/ 246 levels "", "1_reihe", "100",...: 42 21 118 108 105 236 118 118 44 15 ...
 $ odometer        : Factor w/ 13 levels "10,000km","100,000km",...: 4 4 11 11 4 4 4 4 4 4 ...
 $ monthOfRegistration: int  3 6 7 6 7 4 8 12 10 7 ...
 $ fuelType        : Factor w/ 8 levels "", "andere", "benzin",...: 8 3 3 3 3 5 3 5 1 3 ...
 $ brand           : Factor w/ 40 levels "alfa_romeo","audi",...: 26 3 39 33 11 5 39 39 31 28 ...
 $ notRepairedDamage : Factor w/ 3 levels "", "ja", "nein": 3 3 3 3 3 1 1 3 3 3 ...
 $ dateCreated      : Factor w/ 76 levels "2015-06-11 00:00:00",...: 64 73 64 50 70 59 58 54 60 54 ...
 $ nrOfPictures     : int  0 0 0 0 0 0 0 0 0 0 ...
 $ postalCode       : int  79588 71034 35394 33729 39218 22962 31535 53474 7426 15749 ...
 $ lastSeen         : Factor w/ 39481 levels "2016-03-05 14:45:46",...: 31336 33321 34787 4785 21264 3 ...

```

Many of these variables are factor variables though they don't need to. Such variables include `lastSeen`, `dateCreated`, `price`, and `odometer`.

We can check if there are NAs, too. Thankfully none of the variables are null.

```
sapply(auto_info, function(x) sum(is.na(x)))
```

```

      dateCrawled      name      seller
           0           0           0
      offerType      price      abtest
           0           0           0
      vehicleType yearOfRegistration      gearbox
           0           0           0
      powerPS      model      odometer
           0           0           0
monthOfRegistration      fuelType      brand

```

0	0	0
notRepairedDamage	dateCreated	nrOfPictures
0	0	0
postalCode	lastSeen	
0	0	

We can get a five-number summary of this data.

```
summary(auto_info)
```

```

      dateCrawled      name
2016-03-05 16:57:05:   3 Ford_Fiesta      : 78
2016-03-08 10:40:35:   3 BMW_316i         : 75
2016-03-09 11:54:38:   3 Volkswagen_Golf_1.4: 75
2016-03-10 15:36:24:   3 BMW_318i         : 72
2016-03-11 22:38:16:   3 Volkswagen_Polo   : 72
2016-03-12 16:06:22:   3 BMW_320i         : 71
(Other)      :49982 (Other)      :49557

      seller      offerType      price      abtest
gewerblich:   1 Angebot:49999 $0      : 1421 control:24244
privat      :49999 Gesuch :   1 $500   : 781 test   :25756
              $1,500 : 734
              $2,500 : 643
              $1,000 : 639
              $1,200 : 639
              (Other):45143

      vehicleType      yearOfRegistration      gearbox      powerPS
limousine :12859 Min. :1000      : 2680 Min. : 0.0
kleinwagen:10822 1st Qu.:1999      automatik:10327 1st Qu.: 70.0
kombi      : 9127 Median :2003      manuell :36993 Median : 105.0
              : 5095 Mean  :2005      Mean  : 116.4
bus        : 4093 3rd Qu.:2008      3rd Qu.: 150.0
cabrio     : 3061 Max.  :9999      Max.  :17700.0
(Other)    : 4943

      model      odometer      monthOfRegistration      fuelType
golf : 4024 150,000km:32424 Min. : 0.000 benzin :30107
andere : 3528 125,000km: 5170 1st Qu.: 3.000 diesel :14567
3er : 2761 100,000km: 2169 Median : 6.000      : 4482
      : 2758 90,000km : 1757 Mean  : 5.723 lpg : 691
polo : 1757 80,000km : 1436 3rd Qu.: 9.000 cng : 75
corsa : 1735 70,000km : 1230 Max. :12.000 hybrid : 37
(Other):33437 (Other) : 5814      (Other): 41

      brand      notRepairedDamage      dateCreated
volkswagen :10687 : 9829 2016-04-03 00:00:00: 1946
opel : 5461 ja : 4939 2016-03-20 00:00:00: 1893
bmw : 5429 nein:35232 2016-03-21 00:00:00: 1886
mercedes_benz: 4734 2016-04-04 00:00:00: 1844
audi : 4283 2016-03-12 00:00:00: 1831
ford : 3479 2016-03-14 00:00:00: 1761
(Other) :15927 (Other) :38839

      nrOfPictures      postalCode      lastSeen
Min. :0 Min. : 1067 2016-04-07 06:17:27: 8
1st Qu.:0 1st Qu.:30451 2016-04-06 06:17:24: 7
Median :0 Median :49577 2016-04-06 21:17:51: 7
Mean :0 Mean :50814 2016-04-07 03:16:17: 7

```

```

3rd Qu.:0      3rd Qu.:71540  2016-04-05 16:44:47: 6
Max. :0      Max. :99998  2016-04-06 01:16:01: 6
                                (Other) :49959

```

Taking a brief look, the number of pictures (`nrOfPictures`) is completely zero, so this column can be removed.

```

# get the column number of that variable
pictures_col <- grep("Pictures", colnames(auto_info))
# remove the column
auto_info <- auto_info[, -pictures_col]

```

Cleaning the data structure

Earlier it was mentioned some of the variables have improper datatypes. Some of these variables, like `price` and `odometer` have extra characters (\$) and km). Even if no analysis is to be done on them, it's still helpful to reformat them anyway just in case.

Here's a rundown of what to transform each variable's datatype in to.

```

var_transform <- data.frame(
  'Variable Name' = c('dateCrawled', 'name',
                      'price', 'model', 'odometer',
                      'brand', 'dateCreated', 'lastSeen'),

  'Convert To' = c('Date', 'String', 'Numeric',
                   'String', 'Integer', 'String', 'Date', 'Date'))

var_transform %>%
  `colnames<-`(c('Variable Name', 'Convert To')) %>%
  kable(align = rep('c', 2)) %>%
  kable_styling(bootstrap_options = "striped", full_width = F)

```

Variable Name	Convert To
dateCrawled	Date
name	String
price	Numeric
model	String
odometer	Integer
brand	String
dateCreated	Date
lastSeen	Date

Let's do that now in the order of the table above.

```

auto_info$dateCrawled <- ymd_hms(auto_info$dateCrawled) %>% date()
auto_info$name <- as.character(auto_info$name)
# remove "$" and any commas from the price column
auto_info$price %<>% as.character() %>% gsub(",", "", .) %>%
  sub("\\$", "", .) %>% as.numeric()
auto_info$model <- as.character(auto_info$model)
# remove "km" and "," and make values numeric in the odometer column
auto_info$odometer %<>% as.character() %>% sub("km", "", .) %>%
  sub(",", "", .) %>% as.numeric()
auto_info$brand <- as.character(auto_info$brand)

```

```
auto_info$dateCreated <- ymd_hms(auto_info$dateCreated) %>% date()
auto_info$lastSeen <- ymd_hms(auto_info$lastSeen) %>% date()
```

Exploring Odometer and Price

We'll examine odometer and price for any patterns, beginning with odometer.

```
auto_info %>%
  group_by(`Odometer Value (km)` = odometer) %>%
  summarise(Total = length(`Odometer Value (km)`)) %>%
  arrange(desc(`Odometer Value (km)`)) %>%
  kable(align = rep('c',2)) %>%
  kable_styling(bootstrap_options = "striped", full_width = F)
```

Odometer Value (km)	Total
150000	32424
125000	5170
100000	2169
90000	1757
80000	1436
70000	1230
60000	1164
50000	1027
40000	819
30000	789
20000	784
10000	264
5000	967

Clearly, the majority of used cars have traveled farther.

Let's look at price.

```
price_summary <- auto_info %>% select(price) %>%
  group_by(`Price($)` = price) %>%
  summarise(`Total Cars` = length(`Price($)`))

price_summary %>% arrange(`Price($)` ) %>% head() %>%
  kable(align = rep('c',2)) %>%
  kable_styling(bootstrap_options = "striped", full_width = F)
```

Price(\$)	Total Cars
0	1421
1	156
2	3
3	1
5	2
8	1

```
price_summary %>% arrange(desc(`Price($)`)) %>%
  head() %>% kable(align = rep('c',2)) %>%
  kable_styling(bootstrap_options = "striped", full_width = F)
```

Price(\$)	Total Cars
99999999	1
27322222	1
12345678	3
11111111	2
10000000	1
3890000	1

Amazingly, some cars are listed as \$0, though that represents only 2% of the cars. The most expensive car is \$99,999,999!

Exploring Ad Dates

The `lastSeen` column records the date the web crawler last saw any listing, which allows us to determine on what day a listing was removed, presumably because the car was sold. Let's take a look and see if any patterns emerge.

```
ad_timeline <- auto_info %>%
  group_by(`Last Seen` = lastSeen) %>%
  summarise(Total = length(`Last Seen`))


ad_timeline %>% head() %>%
  kable(aligned = rep('c', 2)) %>%
  kable_styling(bootstrap_options = "striped", full_width = F)
```

Last Seen	Total
2016-03-05	54
2016-03-06	221
2016-03-07	268
2016-03-08	380
2016-03-09	493
2016-03-10	538

Better to make a time series plot.

```
ggplot(ad_timeline) +
  geom_line(aes(`Last Seen`, Total), col = 'steel blue') +
  xlab("Day") + ylab("Total Ads") +
  ggtitle("Car Listings Removed Between 3/5/16 and 4/7/16") +
  theme(plot.title = element_text(hjust = 0.5))
```

Car Listings Removed Between 3/5/16 and 4/7/16



The number of car listings taken down is roughly uniform until April 4 (the last three days of the dataset). It's unclear what's behind this behavior but further research can be done to determine the cause of (an apparent) buying frenzy.

Fixing Incorrect Registration Year Data

Looking back at the `summary(auto_info)` output, you'll see the minimum value for the `yearOfRegistration` column is 1000 and the maximum value is 9999, obviously incorrect data.

```
summary(auto_info$yearOfRegistration)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
1000	1999	2003	2005	2008	9999

Furthermore, there is mismatch between the year an eBay ad was posted and the year a car was first registered. In other words, it's not possible for a car to be first registered after the listing was first seen. However, this anomaly exists.

```
# find latest year for car registration
auto_info %>% select(yearOfRegistration) %>%
  filter(yearOfRegistration < 2020) %>%
  arrange(desc(yearOfRegistration)) %>%
  head(1)
```

```
yearOfRegistration
1                2019
```

```
# find latest year a car listing was made
year(auto_info$dateCreated) %>% max()
```

```
[1] 2016
```

So the latest listing is from 2016 but the latest car registration year is 2019. For simplicity we'll assume the earliest valid registration is somewhere in the early 20th century. We can count the number cars with registration outside 1900-2016 and see if those rows can safely be removed, or if more custom logic is needed.

```
auto_info %>%  
  filter(!between(yearOfRegistration, 1900, 2016)) %>%  
  nrow()/nrow(auto_info)
```

```
[1] 0.03944
```

Car registrations outside the 1900-2016 range account for less than 4% of the complete dataset, so these can be safely removed.

```
auto_info %<>% filter(between(yearOfRegistration, 1900, 2016))
```

Exploring Price by Brand

Let's first look at the number of car listings by brand.

```
top_listings <- auto_info %>%  
  select(brand) %>% group_by(Brand = brand) %>%  
  summarise(Total = length(Brand)) %>%  
  arrange(desc(Total))
```

```
top_listings
```

```
# A tibble: 40 x 2  
  Brand      Total  
  <chr>      <int>  
1 volkswagen 10188  
2 bmw        5284  
3 opel       5195  
4 mercedes_benz 4580  
5 audi       4149  
6 ford       3352  
7 renault    2274  
8 peugeot    1418  
9 fiat       1242  
10 seat       873  
# ... with 30 more rows
```

As one might guess (though it need not be the true), the highest proportion of cars are German in origin, representing four out of the top 5 brands. The price analysis will focus on the top five brands.

```
top_five_listings = head(top_listings$Brand, 5)
```

```
auto_info %>%  
  filter(brand %in% top_five_listings) %>%  
  group_by(Brand = brand) %>%  
  summarise(`Mean Price` = mean(price))
```

```
# A tibble: 5 x 2  
  Brand      `Mean Price`  
  <chr>      <dbl>  
1 audi       9094.
```



```
2 bmw 8335.
3 mercedes_benz 30317.
4 opel 5253.
5 volkswagen 6516.
```

Mercedes has the highest mean price, but is there an outlier? Yes.

```
top_n(auto_info, 1, auto_info$price) %>% select(price, brand)
```

```
price brand
1 1e+08 mercedes_benz
```

The highest priced car (at \$99,999,999) belongs to this brand. So why is 1e08 (\$100,000,000) shown? This is the well known problem of precision in computer science. Just to prove I'm not bluffing, take a look at this:

```
x = 99999999
x
```

```
[1] 1e+08
```

It's outside the scope of this report but you can find more info (if not familiar already with the subject) beginning [here](#).

If we remove this row, let's see how the mean prices change.

```
# remove outlier
auto_info <- auto_info[-which(auto_info$price == max(auto_info$price)), ]
# find mean prices
auto_info %>% filter(brand %in% top_five_listings) %>%
  group_by(Brand = brand) %>%
  summarise(`Mean Price` = mean(price))
```

```
# A tibble: 5 x 2
  Brand      `Mean Price`
  <chr>          <dbl>
1 audi          9094.
2 bmw           8335.
3 mercedes_benz 8485.
4 opel          5253.
5 volkswagen    6516.
```

Now it's Audi that features the highest mean price.

Exploring Mileage

The final variable of exploration is mileage. Which brands have higher mileage listed, and does this correlate with price?

```
auto_info %>% filter(brand %in% top_five_listings) %>%
  group_by(Brand = brand) %>%
  select(Brand, price, odometer) %>%
  summarise(`Mean Price` = mean(price), `Mean Mileage` = mean(odometer)) %>%
  arrange(desc(`Mean Mileage`))
```

```
# A tibble: 5 x 3
  Brand      `Mean Price` `Mean Mileage`
  <chr>          <dbl>          <dbl>
1 bmw           8335.         132435.
2 mercedes_benz 8485.         130856.
```

3 audi	9094.	129288.
4 opel	5253.	129227.
5 volkswagen	6516.	128730.

There isn't much variability in mean mileage, but cars with higher mileage tend to be more expensive.

Conclusion

This project explores used car data from Germany's eBay website. The analysis covered questions regarding:

- data structure
- most popular brands
- outliers
- price
- mileage
- car listing dates